

In [117...

```
import pandas as pd
import numpy as np
```

In [118...

```
df_train = pd.read_csv("train.csv")
df_test = pd.read_csv("test.csv")
```

In [119...

```
# Check for null values in training as well as testing data
```

In [120...

```
df_train.isnull().sum()
```

Out[120...

```
ID      0
y        0
X0       0
X1       0
X2       0
..
X380     0
X382     0
X383     0
X384     0
X385     0
Length: 378, dtype: int64
```

In [121...

```
df_test.isnull().sum()
```

Out[121...

```
ID      0
X0       0
X1       0
X2       0
X3       0
..
X380     0
X382     0
X383     0
X384     0
X385     0
Length: 377, dtype: int64
```

In [122...

```
df_train.head()
```

Out[122...

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X382
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	1
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0

5 rows × 378 columns



In [123...

```
df_train.dtypes.value_counts()
```

Out[123... int64 369
object 8
float64 1
dtype: int64

In [124... df_test.dtypes.value_counts()

Out[124... int64 369
object 8
dtype: int64

In [125... df_test.head()

Out[125...

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	0
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	0
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	0
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	0
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	0

5 rows × 377 columns



In [126... *# Drop the ID column from the tet data and ID and y from the training data*
No null values are found in the data
Check for the columns with zero variance and remove those columns
Before that seperate the dependent and independent variable
y_train = df_train["y"]
df_id = df_train["ID"]
df_train = df_train.drop(["ID", "y"], axis = 1)
df_test = df_test.drop("ID", axis = 1)

In [127... df_train.head()

Out[127...

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	...	X375	X376	X377	X378	X379	X380	X382
0	k	v	at	a	d	u	j	o	0	0	...	0	0	1	0	0	0	0
1	k	t	av	e	d	y	l	o	0	0	...	1	0	0	0	0	0	0
2	az	w	n	c	d	x	j	x	0	0	...	0	0	0	0	0	0	1
3	az	t	n	f	d	x	l	e	0	0	...	0	0	0	0	0	0	0
4	az	v	n	f	d	h	d	n	0	0	...	0	0	0	0	0	0	0

5 rows × 376 columns



In [128... df_test.head()

Out[128...

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	...	X375	X376	X377	X378	X379	X380	X382
--	----	----	----	----	----	----	----	----	-----	-----	-----	------	------	------	------	------	------	------

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	...	X375	X376	X377	X378	X379	X380	X382
0	az	v	n	f	d	t	a	w	0	0	...	0	0	0	1	0	0	0
1	t	b	ai	a	d	b	g	y	0	0	...	0	0	1	0	0	0	0
2	az	v	as	f	d	a	j	j	0	0	...	0	0	0	1	0	0	0
3	az	l	n	f	d	z	l	n	0	0	...	0	0	0	1	0	0	0
4	w	s	as	c	d	y	i	m	0	0	...	1	0	0	0	0	0	0

5 rows × 376 columns

In [129...
y_train.head()

Out[129... 0 130.81
1 88.53
2 76.26
3 80.62
4 78.02
Name: y, dtype: float64

In [130... *# Seperate the categorical and numerical data for the training and the testing data*
So that we can easly preprocess the data
df_num_train = df_train.select_dtypes(exclude = np.object)
df_cat_train = df_train.select_dtypes(include = np.object)
df_num_test = df_test.select_dtypes(exclude = np.object)
df_cat_test = df_test.select_dtypes(include = np.object)

<ipython-input-130-7c8aa6cfa58a>:3: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

df_num_train = df_train.select_dtypes(exclude = np.object)
<ipython-input-130-7c8aa6cfa58a>:4: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

df_cat_train = df_train.select_dtypes(include = np.object)
<ipython-input-130-7c8aa6cfa58a>:5: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

df_num_test = df_test.select_dtypes(exclude = np.object)
<ipython-input-130-7c8aa6cfa58a>:6: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>
df_cat_test = df_test.select_dtypes(include = np.object)

In [131...
df_num_train.head()

Out[131... X10 X11 X12 X13 X14 X15 X16 X17 X18 X19 ... X375 X376 X377 X378 X379 X382

0	0	0	0	1	0	0	0	0	1	0	...	0	0	1	0	0
1	0	0	0	0	0	0	0	0	1	0	...	1	0	0	0	0

	X10	X11	X12	X13	X14	X15	X16	X17	X18	X19	...	X375	X376	X377	X378	X379	X380
2	0	0	0	0	0	0	0	1	0	0	...	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0

5 rows × 368 columns



In [132...

df_num_test.head()

Out[132...

	X10	X11	X12	X13	X14	X15	X16	X17	X18	X19	...	X375	X376	X377	X378	X379	X380
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	1	0	0
1	0	0	0	0	0	0	0	0	0	1	...	0	0	1	0	0	0
2	0	0	0	0	1	0	0	0	0	0	...	0	0	0	1	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	1	0	0
4	0	0	0	0	1	0	0	0	0	0	...	1	0	0	0	0	0

5 rows × 368 columns



In [133...

```
# Now search for the columns which has zero variance and remove them
# Because those columns are of no use for us to build a model
v = np.array(df_num_train.var())
col_names = list(df_num_train.columns)
for i in range(368):
    if v[i] ==0:
        df_num_train = df_num_train.drop(col_names[i], axis =1)
        df_num_test = df_num_test.drop(col_names[i], axis = 1)
```

In [134...

```
## All the columns with zero variance is reoved from the taining and testing numeric
df_num_train.head()
```

Out[134...

	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	...	X375	X376	X377	X378	X379	X380
0	0	0	1	0	0	0	0	1	0	0	...	0	0	1	0	0	0
1	0	0	0	0	0	0	0	1	0	0	...	1	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0	...	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0

5 rows × 356 columns



In [135...

df_num_test.head()

Out[135...

	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	...	X375	X376	X377	X378	X379	X380
--	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	------	------	------	------	------	------

	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	...	X375	X376	X377	X378	X379	X380
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	1	0	0
1	0	0	0	0	0	0	0	0	1	0	...	0	0	1	0	0	0
2	0	0	0	1	0	0	0	0	0	0	...	0	0	0	1	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	1	0	0
4	0	0	0	1	0	0	0	0	0	0	...	1	0	0	0	0	0

5 rows × 356 columns



In [136... *# Now apply the label encoder on the categorical columns in the training and testing*
from sklearn.preprocessing *import* LabelEncoder

In [137...
 le = LabelEncoder()
 col_names = list(df_cat_train.columns)
for i *in* col_names:
 df_cat_train[i] = le.fit_transform(df_cat_train[i])
 df_cat_test[i] = le.fit_transform(df_cat_test[i])

<ipython-input-137-57cc79ed0218>:4: SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

df_cat_train[i] = le.fit_transform(df_cat_train[i])
 <ipython-input-137-57cc79ed0218>:5: SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 df_cat_test[i] = le.fit_transform(df_cat_test[i])

In [138... df_cat_train.head()

Out[138...

	X0	X1	X2	X3	X4	X5	X6	X8
0	32	23	17	0	3	24	9	14
1	32	21	19	4	3	28	11	14
2	20	24	34	2	3	27	9	23
3	20	21	34	5	3	27	11	4
4	20	23	34	5	3	12	3	13

In [139... df_cat_test.head()

Out[139...

	X0	X1	X2	X3	X4	X5	X6	X8
0	21	23	34	5	3	26	0	22
1	42	3	8	0	3	9	6	24

	X0	X1	X2	X3	X4	X5	X6	X8
2	21	23	17	5	3	0	9	9
3	21	13	34	5	3	31	11	13
4	45	20	17	2	3	30	8	12

In [140...

df_cat_train.dtypes.value_counts()

Out[140...

int32 8
dtype: int64

In [141...

df_cat_test.dtypes.value_counts()

Out[141...

int32 8
dtype: int64

In [142...

```
# Now apply the MinMaxScaler technique on the categorical data sets
from sklearn.preprocessing import MinMaxScaler
```

In [143...

```
mn = MinMaxScaler()
df_cat_train_1 = mn.fit_transform(df_cat_train)
df_cat_test_1 = mn.fit_transform(df_cat_test)
df_cat_train_sc = pd.DataFrame(df_cat_train_1, index = df_cat_train.index, columns =
df_cat_test_sc = pd.DataFrame(df_cat_test_1, index = df_cat_test.index, columns = df
```

In [144...

```
## All the categorical data is chnaged into the numerical data
# Now concat the numerical and categorical data set of the training and testing data
df_final_train = pd.concat([df_num_train, df_cat_train_sc], axis = 1)
df_final_test = pd.concat([df_num_test, df_cat_test_sc], axis = 1)
```

In [145...

df_final_train.head()

Out[145...

	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	...	X384	X385	X0	X1
0	0	0	1	0	0	0	0	1	0	0	...	0	0	0.695652	0.884615
1	0	0	0	0	0	0	0	1	0	0	...	0	0	0.695652	0.807692
2	0	0	0	0	0	0	1	0	0	0	...	0	0	0.434783	0.923077
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0.434783	0.807692
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0.434783	0.884615

5 rows × 364 columns



In [146...

df_final_test.head()

Out[146...

	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	...	X384	X385	X0	X1
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0.4375	0.884615
1	0	0	0	0	0	0	0	0	1	0	...	0	0	0.8750	0.115385

	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	...	X384	X385	X0	X1	
2	0	0	0	1	0	0	0	0	0	0	...	0	0	0.4375	0.884615	0.3863
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0.4375	0.500000	0.7727
4	0	0	0	1	0	0	0	0	0	0	...	0	0	0.9375	0.769231	0.3863

5 rows × 364 columns



In [147... *## Now apply the Dimensionality Reduction via using PCA*
from sklearn.decomposition **import** PCA

In [148...
 n_comp = 12
 pca = PCA(n_components = n_comp, random_state = 420)
 df_train_pca = pca.fit_transform(df_final_train)
 df_test_pca = pca.fit_transform(df_final_test)

In [149... *## Number of columns are reduced to 12 for both training as well as the testing data*
But the data is in nd array format so covert that into data frame
 df_train_pca2 = pd.DataFrame(df_train_pca, columns = ["PC1", "PC2", "PC3", "PC4", "PC5", "PC6", "PC7", "PC8", "PC9", "PC10", "PC11", "PC12"])
 df_test_pca2 = pd.DataFrame(df_test_pca, columns = ["PC1", "PC2", "PC3", "PC4", "PC5", "PC6", "PC7", "PC8", "PC9", "PC10", "PC11", "PC12"])

In [150... df_train_pca2.head()

Out[150...

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	
0	0.682272	2.217390	1.233625	0.885738	1.401422	0.054223	0.654832	-0.937322	0.192252	-0
1	-0.279051	1.164201	-0.764263	-0.660639	0.237863	0.066804	1.237285	-0.530337	-0.108870	0
2	-1.018083	2.979512	0.558557	2.540751	-0.926714	3.282631	-0.940264	0.557082	-0.925952	-0
3	-0.658559	2.545045	-0.425408	2.997377	-1.681632	3.134975	0.074145	0.084039	-1.072854	0
4	-0.652313	2.370739	-0.583703	3.194208	-1.999394	3.167652	-0.143351	0.229232	-1.754659	-0



In [151... df_test_pca2.head()

Out[151...

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	
0	-0.425629	0.308747	-2.524555	2.641045	-1.995091	3.513185	-0.302433	0.095607	-1.952755	-0
1	3.714128	1.677736	-0.003393	-0.604436	1.448695	-0.106079	-1.490445	-1.305825	-1.211747	0
2	-1.204703	0.513145	-0.911797	1.237162	-0.392674	3.213377	-0.737050	0.718334	-1.504446	-0
3	-0.367696	0.181923	-2.419289	2.552553	-2.111593	3.544317	-0.253927	0.027803	-1.974823	-0
4	-2.860633	0.470359	0.775156	-1.612700	0.796252	0.178710	-0.133292	0.098793	-0.079083	-0



In [152... *## Now all preprocessing steps are done*
Now apply the train_test_split to seperate the training and validation data

```
from sklearn.model_selection import train_test_split
```

```
In [153...  ## Now define the dependent and independent variables
x = df_train_pca2
y = y_train
```

```
In [154... X_train, X_valid, Y_train, Y_valid = train_test_split(x, y, test_size = 0.2, random_
```

```
In [155...  ## Training and validation data is created
# Now train the model using XGBoost algorithm
import xgboost as xgb
```

```
In [156... d_train = xgb.DMatrix(X_train, label = Y_train)
d_valid = xgb.DMatrix(X_valid, label = Y_valid)
d_test = xgb.DMatrix(df_test_pca2)
```

```
In [157... params = {}
params["objective"] = "reg:linear"
params["eta"] = 0.02
params["max_depth"] = 4
```

```
In [158... w = [(d_train, "train"), (d_valid, "valid")]
```

```
In [159...  ## import the r2_score metric for the evaluation of the model
from sklearn.metrics import r2_score
```

```
In [160...  ## create a function to
#generate the r2_score
def score(preds, dtrain):
    labels = dtrain.get_label()
    return "r2", r2_score(labels, preds)
```

```
In [161...  ## Now lets train the model
trained = xgb.train(params, d_train, 1000, w, early_stopping_rounds = 50, feval = sc
```

[13:05:50] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/objective/regression_obj.cu:171: reg:linear is now deprecated in favor of reg:squarederror.

[0]	train-rmse:99.14659	train-r2:-58.35085	valid-rmse:98.26758	valid-r2:-67.64398
[10]	train-rmse:81.25641	train-r2:-38.86453	valid-rmse:80.39062	valid-r2:-44.94020
[20]	train-rmse:66.67461	train-r2:-25.84061	valid-rmse:65.79161	valid-r2:-29.76968
[30]	train-rmse:54.79623	train-r2:-17.12896	valid-rmse:53.86914	valid-r2:-19.62824
[40]	train-rmse:45.13720	train-r2:-11.30101	valid-rmse:44.16470	valid-r2:-12.86541
[50]	train-rmse:37.29795	train-r2:-7.39927	valid-rmse:36.29005	valid-r2:-8.36175
[60]	train-rmse:30.94680	train-r2:-4.78233	valid-rmse:29.90724	valid-r2:-5.35821
[70]	train-rmse:25.83114	train-r2:-3.02864	valid-rmse:24.75159	valid-r2:-3.35500

[80]	train-rmse:21.73255	train-r2:-1.85163	valid-rmse:20.61720	vali
	d-r2:-2.02163			
[90]	train-rmse:18.47799	train-r2:-1.06149	valid-rmse:17.33865	vali
	d-r2:-1.13704			
[100]	train-rmse:15.91879	train-r2:-0.53000	valid-rmse:14.75732	vali
	d-r2:-0.54809			
[110]	train-rmse:13.93206	train-r2:-0.17193	valid-rmse:12.76408	vali
	d-r2:-0.15814			
[120]	train-rmse:12.40632	train-r2:0.07070	valid-rmse:11.25407	vali
	d-r2:0.09967			
[130]	train-rmse:11.25278	train-r2:0.23548	valid-rmse:10.13224	vali
	d-r2:0.27022			
[140]	train-rmse:10.39089	train-r2:0.34811	valid-rmse:9.30835	vali
	d-r2:0.38407			
[150]	train-rmse:9.75497	train-r2:0.42546	valid-rmse:8.72695	vali
	d-r2:0.45861			
[160]	train-rmse:9.29393	train-r2:0.47848	valid-rmse:8.32602	vali
	d-r2:0.50721			
[170]	train-rmse:8.97829	train-r2:0.51330	valid-rmse:8.06254	vali
	d-r2:0.53791			
[180]	train-rmse:8.75464	train-r2:0.53725	valid-rmse:7.89498	vali
	d-r2:0.55692			
[190]	train-rmse:8.58798	train-r2:0.55470	valid-rmse:7.78532	vali
	d-r2:0.56914			
[200]	train-rmse:8.46478	train-r2:0.56738	valid-rmse:7.72132	vali
	d-r2:0.57620			
[210]	train-rmse:8.36830	train-r2:0.57719	valid-rmse:7.68090	vali
	d-r2:0.58062			
[220]	train-rmse:8.29351	train-r2:0.58471	valid-rmse:7.66025	vali
	d-r2:0.58287			
[230]	train-rmse:8.22629	train-r2:0.59142	valid-rmse:7.64971	vali
	d-r2:0.58402			
[240]	train-rmse:8.17027	train-r2:0.59696	valid-rmse:7.64299	vali
	d-r2:0.58475			
[250]	train-rmse:8.12457	train-r2:0.60146	valid-rmse:7.64451	vali
	d-r2:0.58459			
[260]	train-rmse:8.08198	train-r2:0.60563	valid-rmse:7.65097	vali
	d-r2:0.58388			
[270]	train-rmse:8.05113	train-r2:0.60863	valid-rmse:7.65869	vali
	d-r2:0.58304			
[280]	train-rmse:8.00920	train-r2:0.61270	valid-rmse:7.66112	vali
	d-r2:0.58278			
[290]	train-rmse:7.96914	train-r2:0.61656	valid-rmse:7.66797	vali
	d-r2:0.58203			
[292]	train-rmse:7.96561	train-r2:0.61690	valid-rmse:7.66941	vali
	d-r2:0.58188			

In [162... *## Now predict the test values using xgboost*

```
predict = trained.predict(d_test)
```

In [168... *## Now create a seperate data frame
so that we can clearly see the results*

```
res = pd.DataFrame()
res["ID"] = df_id
res["y"] = predict
res.head()
```

Out[168... **ID** **y**

0	0	78.084824
1	6	92.252342
2	7	81.842590
3	9	78.531166

	ID	y
4	13	108.980766

In [172...

```
## Now save the result in a csv file  
res.to_csv("Final_Output.csv", index = False)
```